



8.Transaction

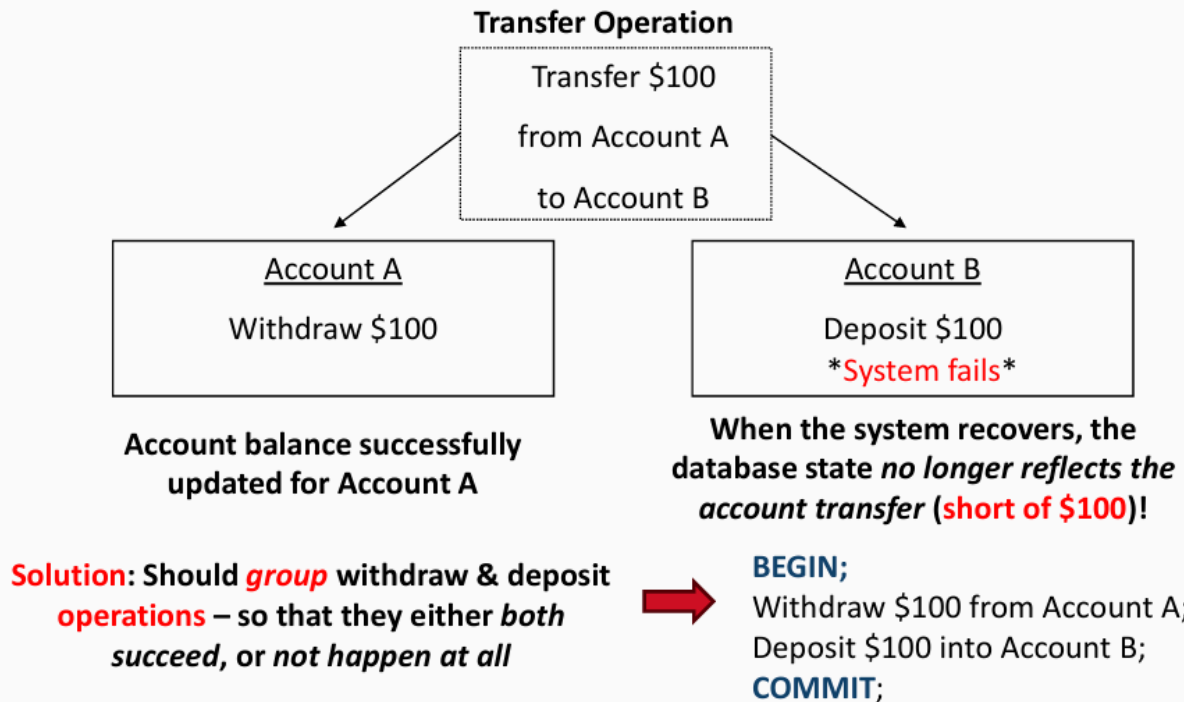
8.1 ACID

- Atomicity : transaction is performed entirely or not performed at all.
- Consistency : A correct execution of a transaction must take a database from one consistent state to another.
- Isolation : Effect of multiple transaction is the same as the transaction running one after another(并发 concurrent)
- Durability : Once a transaction changes the database and the changes are committed, these changes should never be lost.

8.2 Transaction

8.2.1 Example of no Transaction

What could happen if we *did not have a transaction*?



8.2.2 Definition

A program consisting **one or more** SQL statements.

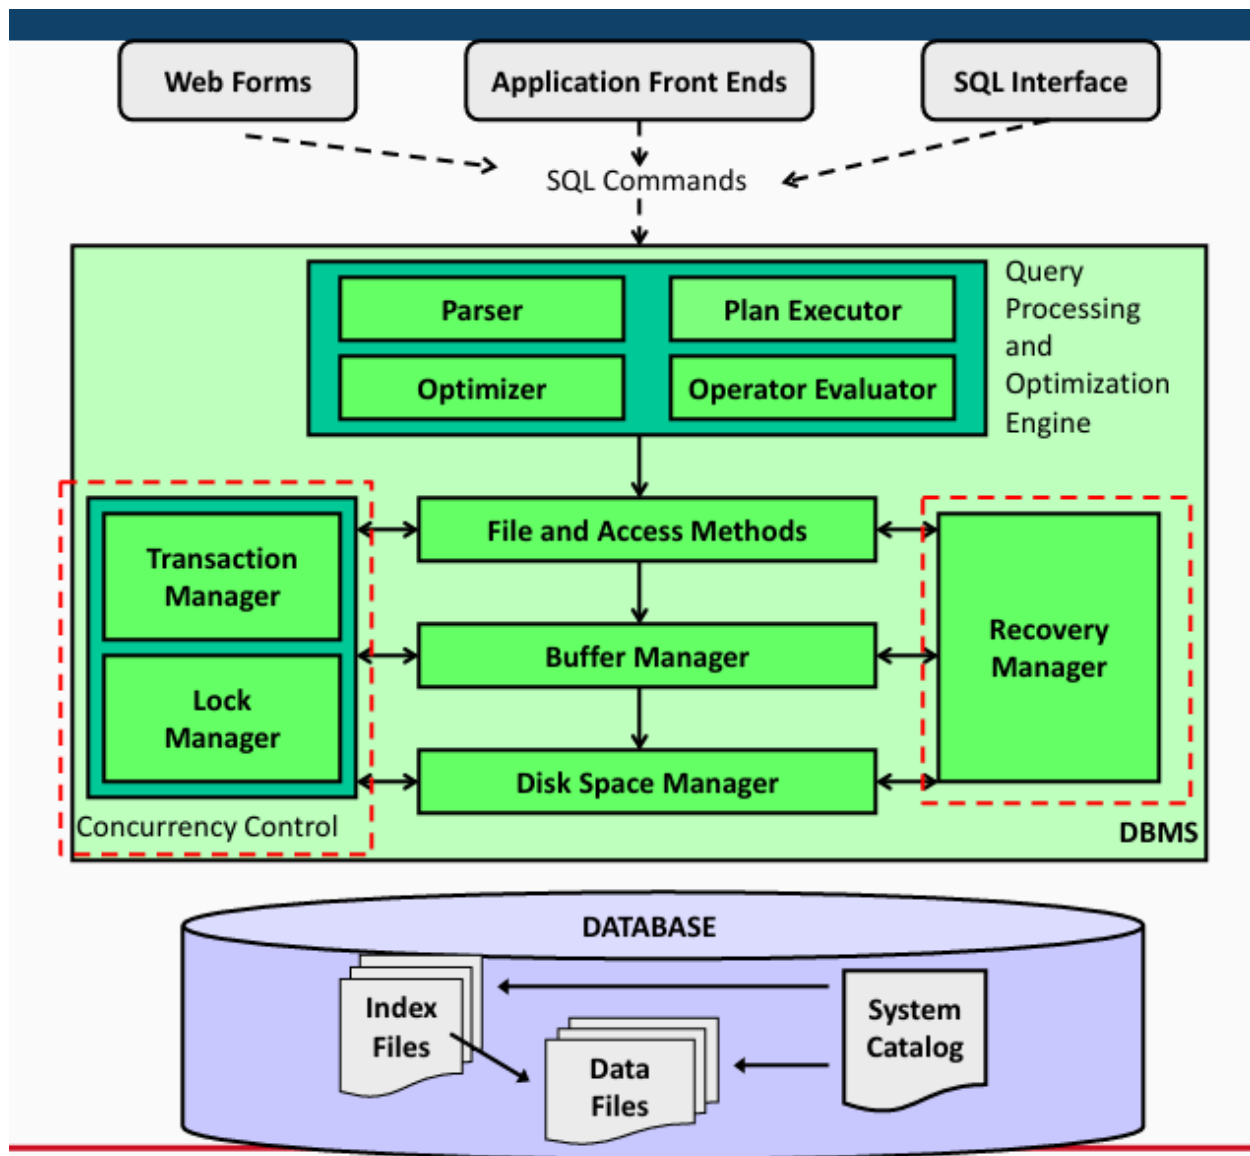
- Executed as an atomic unit
 - atomicity means executes fully or not at all

Another way to describe a transaction.

- A program that is executed to change the database state in a correct way.

A collection of one or more operations which reflect a discrete unit of work.

8.2.3 Internal Structure of a DBMS



8.2.4 Transaction consistency

- Before transaction
 - Assuming the database is in consistent state (satisfy all constraints).
- After transaction
 - All database constraints are satisfied.
 - The new database state satisfies the specification of the transaction.

Consistency means any transaction can only modify data in allowed way, which must agree with all defined rules, constraints and triggers

- Example of a *bad transaction* for a bank transfer:
 - Withdraw \$100 from account A, but only deposit \$90 into account B.

We can select when to enforce the database consistency in the transaction.

Example: We may *defer* the enforcement of integrity constraints.

```
CREATE TABLE UnitOfStudy (  
  uos_code          VARCHAR(8),  
  title             VARCHAR(20),  
  lecturer_id       INTEGER,  
  credit_points     INTEGER,  
  
  CONSTRAINT UnitOfStudy_PK PRIMARY KEY (uos_code),  
  CONSTRAINT UnitOfStudy_FK FOREIGN KEY (lecturer_id)  
    REFERENCES Lecturer DEFERRABLE INITIALLY IMMEDIATE  
);  
  
BEGIN;  
  SET CONSTRAINTS UnitOfStudy_FK DEFERRED;  
  INSERT INTO UnitOfStudy VALUES('INFO1000', 'Graphics', 3, 6);  
  INSERT INTO Lecturer VALUES(3, 'Steve', CSE);  
  SET CONSTRAINTS UnitOfStudy_FK IMMEDIATE ;  
COMMIT;
```

UnitOfStudy			
<u>uos_code</u>	title	lecturer_id	credit_points
COMP9120	DBMS	1	6
COMP9007	Algorithm	2	6

Lecturer		
<u>Lecturer_id</u>	name	department
1	Adam	CSE
2	Lily	IT

does not exist yet!

8.2.5 Transaction atomic

Logs all actions that would need to be undone if the transaction is aborted (incomplete).

- all actions undone
- do a redo

- If the transaction successfully completes, it is said to have committed.
- else, it is said to have been aborted

8.2.6 Durable

Once the transaction is committed, its effects should persist in a database.

- Transaction logs
 - use stable storage as a log to store a history of modifications made to the database.
 - Recovery Manager should take responsible to the transaction logs.
 - Every transaction has a log associated with it.
 - If a transaction aborts, depending on the recovery protocol, use the log to undo(rollback)/redo(copy the log value of an item from stable storage to disk) the transaction.

8.2.7 Isolation

Transaction can run concurrently, their operations can be interleaved (交错) .

The interleaving is usually decided by host operating system.

If no intervention from transaction manager, it may leave incorrect state.

So the **Control concurrent** is **correctness** but also **efficiency** .

- 3 types of problems
 - Lost update : concurrent update make the former update lost.
 T1 : READ X
 T2 : READ X
 T1 : WRITE X = X + 20

T2 : WRITE X = X + 30

When concurrent happen, return ONLY X + 30

- Temporary update : update an item and then fails, another read the wrong temporary value

T1: Update X to 200

T2: Read X = 200 ← 读取了临时、不合法的值

T1: Rollback ← X 实际上还是 100

- Incorrect summary : when a transaction is updating an aggregate of items, another is reading or updating the value

X = 100 Y = 50

T1: Read X = 100

T1: Write X = 50

T2: Read X = 50 ← 新值

T2: Read Y = 50 ← 旧值

T1 : WRITE Y = X + Y

T2: WRITE SUM = X + Y = 100

→ Total = 100 ← 实际应该是 X=50, Y=100 → Total=150

- 4 isolation levels in SQL standard
 - read uncommitted : dirty read — A transaction can read data written by a concurrent uncommitted transaction.
 - read committed : no dirty read → Lowest level of PostgreSQL
 - repeatable read : In a transaction, it can only see 1 version of data. (which is data committed before transaction began)
 - serializable : highest level of isolation, concurrent executing = some serial execution of these same transactions.

- **read uncommitted** → allows dirty reads
- **read committed** → addresses temporary update problem
- **repeatable read** → addresses incorrect summary problems
- **serializable** → addresses the lost update problem

```

BEGIN;
SET TRANSACTION ISOLATION LEVEL

{ SERIALIZABLE | REPEATABLE READ | READ COMMITTED | READ UNCOMM
-- 下面执行的所有语句都在该隔离级别下运行
SELECT * FROM Account WHERE balance > 1000;

COMMIT;

```

- Transactions should be isolated

Let us consider two transactions that are run *concurrently*

- Transaction T1 is transferring \$100 from account A to account B.
- T2 credits both accounts with a 5% interest payment.

T1: BEGIN	A=A-100,	B=B+100	COMMIT
T2: BEGIN	A=1.05*A,	B=1.05*B	COMMIT

We assume that all transactions commit,
there is no aborted transaction!

- Serial execution:

$A = A - 100$ $B = B + 100$

THEN

$A = A * 1.05$ $B = 1.05 * B$

- Else

$A = A * 1.05$ $B = B * 1.05$

THEN

$A = A - 100$ $B = B + 100$

- Serial Schedule

A schedule in which all transactions are executed from start to finish, without any interleaving.

- However, Interleaving improves the performance.

Serial makes sure that the result of interleaving = running steps by steps

- T1 T2 interleaving = T1 then T2

Serializability even though is good, but it is expensive to check.

- How to check

- › Assume the following schedule:

T1: $A=A-100,$	$B=B+100$
T2:	$A=1.05*A, B=1.05*B$

- › To do this, we need to see the *DBMS's view of a schedule*

T1: R1 (A) , W1 (A) ,	R1 (B) , W1 (B)
T2:	R2 (A) , W2 (A) , R2 (B) , W2 (B)

R : READ

W : WRITE

- Then Conflict serializability

A schedule is conflict serializable if it is conflict equivalent to a serial schedule.

- Conflicts:

- read different item in any order

T1 : READ(A) T2: READ(B)

- read same item in any order

T1: READ(A) T2: READ(A)

- R/W different data items in any order

T1 : READ(A) T2: WRITE(B)

- reading / writing the same data item

- A read of a transaction T1 , then T2 write it not semantically.

Conflict

T1 : READ(A) T2 : WRITE(A)

- The order of 2 writes of 2 transactions does matter. Last value will depend on which write comes last. Conflict

T1 : WRITE(A) T2 : WRITE(A)

Note

With SQL:

SELECT corresponds to a *Read*

INSERT corresponds to a *Write*

DELETE, UPDATE correspond to

a *Read* followed by *Write*

A schedule is conflict serializable if it is conflict equivalent to some serial schedule.

- 2 schedule is conflict equivalent if

- They involve the same set of operations of the same transactions.
- They order every pair of conflicting operations the same way.

- How to check conflict serializability

- Non-conflicting swapping

Swap the non-conflicting operation until it becomes $T1 \rightarrow T2$

Swapping the operation that has non-conflict operation, the resulting schedule is serial $(T1, T2)$, This means the above schedule is conflict-serializable.

- Another way to test conflict

Precedence Graph (Serialization Graph Testing or SGT)



The edge corresponds to one of the following case :

- T1 executes W before T2 execute R
- T1 executes R before T2 W
- T1 W before T2 W
- No read & read , have edge

Check for cycles, if there is any cycle, then the schedule is not conflict serializable.

If the SGT graph is acyclic then do topological sorting. This would determine a linear order consistent to the partial order.

- But it is expensive!

- Conflict Serializability & Serializability

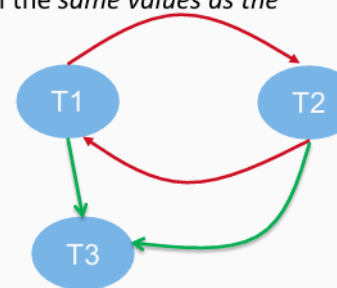
- CS : The given order is identical to a $T1 \rightarrow T2$ (serial) through swapping the non-conflict operation
- S : The result of the given order sequence is identical some of the serial
- CS \rightarrow S but S cannot \rightarrow CS

T1:	W1 (B)		W1 (A)
T2:		W2 (B)	W2 (A)
T3:			W3 (A)

- In this schedule, the final value of A is the value written by T3 and the final value of B is the value written by T2.
- Note that the serial schedule (T1, T2, T3) leaves A and B with the same values as the above schedule. Schedule S is therefore **serializable**.

T1:	W1 (B)	W1 (A)
T2:		W2 (B)
T3:		W3 (A)

- However, schedule S is **not conflict serializable**:
 - It is **not conflict equivalent** to any serial schedule.
- Proof: there is a cycle in the precedence graph.





› Reading Uncommitted Data (“dirty reads”)

T1: R1 (A) , W1 (A) ,	R1 (B) , W1 (B)
T2: R2 (A) , W2 (A) , R2 (B) , W2 (B)	

System Crash

› **Unrepeatable Reads:** may not read same value twice

T1: R (A) ,	R (A)
T2: R (A) , W (A)	

› Overwriting Data produced (written) by another transaction (“**Lost Update**”):

T1: R (A) ,	W (A)
T2: R (A) , W (A)	

多个事务 (Transaction)



构成一个调度 (Schedule)



是否正确? → 判断是否 Serializable



用 Conflict Serializability 来判断



看冲突对 → 画前驱图 → 看是否有环

8.2.8 Lock-based Concurrency Control

A pessimistic protocols : Prevent the conflict happening

- Lock manager maintain a lock table.
- Read locks : “Shared” lock (S)

- Write locks : "Exclusive" lock (X)

T2 Requests \ Held by T1	Shared	Exclusive
Shared	OK	T2 wait on T1
Exclusive	T2 wait on T1	T2 wait on T1

If T1 got a Shared lock, T2 can apply for the shared lock, but applying the Exclusive lock should be blocked.

If T1 got a exclusive lock, T2 cannot apply for any lock.

- The time of release a lock
 - Do it as soon as we have used that item, but may leave database in an inconsistent state.

Example: Consider an airline reservation system where two airline agents are trying to make a booking on the same flight. A schedule could look like this:

T ₁	T ₂
LOCK(X)	
READ(X)	
UNLOCK(X)	
X = X + 1	
	LOCK(X)
	READ(X)
	UNLOCK(X)
	X = X + 1
LOCK(X)	
WRITE(X)	
UNLOCK(X)	
	LOCK(X)
	WRITE(X)
	UNLOCK(X)

This schedule would result in making one single reservation instead of 2!

- Two-Phase Locking

For every transaction

- obtain locks
- perform computations
- release locks and commit

First grant all the locks, release them finally.

- Growing phase : the number of locks can only increase.
- Shrinking phase : the number of locks can only decrease.
- Commit the change to the database

Strict 2PL : all locks are released after commit.

- Deadlock

› Consider the following two transactions:

T1 : R (A) , W (A) , R (B) , W (B)
T2 : R (B) , W (B) , R (A) , W (A)

› A schedule with locks might be:

T1 : S (A) , R (A) , X (A) , W (A) , S (B) ?
T2 : S (B) , R (B) , X (B) , W (B) , S (A) ?

› What is happening here?

- T1 waiting on T2 to release lock on B
- T2 waiting on T1 to release lock on A

DEADLOCK!!

T1 is waiting for the release of T2 while T2 is waiting the release of T1 !

- The time of transaction happening

- T1 holds the lock on A and is requesting a lock on an item B

- T2 holds the lock on item B and requesting a lock on A.

2 way to deal with deadlock :

- Deadlock prevention
 - Static 2-phase locking :
All the transaction will get all the locks or none.
 - Deadlock detection
Use deadlock detection algorithms to detect and aborted the waiting transaction.

each row is an item!

8.2.8 Writing a transaction in SQL

3 key relevant SQL commands

- BEGIN
- COMMIT
- ROLLBACK / ABORT
- also can SET AUTOCOMMIT OFF or SET AUTOCOMMIT ON in pgadmin client
 - auto-commit on each statement is its own transaction

```
BEGIN;  
<你的 SQL 命令>;  
COMMIT;
```

- auto-commit off, you should write begin-commit on you own.