

# COMP9120

Week 5: Database Integrity

Semester 1, 2025

Professor Athman Bouguettaya  
School of Computer Science



THE UNIVERSITY OF  
SYDNEY

# Warming up



THE UNIVERSITY OF  
SYDNEY



# Acknowledgement of Country

*I would like to acknowledge the Traditional Owners of Australia and recognise their continuing connection to land, water and culture. I am currently on the land of the Gadigal people of the Eora nation and pay my respects to their Elders, past, present and emerging.*

---



## **COMMONWEALTH OF AUSTRALIA**

### **Copyright Regulations 1969**

#### **WARNING**

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**).

The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

**Do not remove this notice.**

## › Overview of Integrity Constraints

### › Static Integrity Constraints

- Domain Constraints
- Key / Referential Constraints
- Semantic Integrity Constraints
- Assertions

### › Dynamic Integrity Constraints

- Triggers

## › Integrity Constraint (IC):

A **condition** that must hold **true** for **every instance** of a database:

- ICs are an *integral* part of the initial database *schema design* (through the **create table** command) to ensure the database **integrity** and **consistency**.
- ICs can be *added/updated* at any time (through **alter table table-name add/alter constraint** command).
  - When such a command is executed, the system *first ensures* that the relation *satisfies* the specified *constraint*. If it *does*, the constraint is *added* to the relation; if *not*, the command is *rejected*.
- A **legal** instance of a relation is one that satisfies **all** specified **ICs**.
  - Not *necessarily at all times*.

## › Example of integrity constraints

- Each student ID must be *unique*.
- No two lecturers can have the *same* ID.
- Every school name in the *Unit* relation must have a matching school name in the *School* relation.
- For every student, a *name* must be given.
- The *only* possible grades are either 'F', 'P', 'C', 'D', or 'H'.
- *Valid* lecturer titles are 'Associate Lecturer', 'Lecturer', 'Senior Lecturer', 'Associate Professor', or 'Professor'.
- Students can *only* enrol in the units of study that are *currently* on offer.
- The *sum* of all marks in a course *cannot* at *any time* be *higher* than 100.

› **Why** do we need to **capture integrity constraints**? To ensure:

1. Data is stored in *accordance* with the *real-world meaning* (semantics) of the domain application by:
  - Avoiding *data entry errors* (e.g., inserting a *grade* into the **Student** table which *does not exist*).
2. *Data consistency* (e.g., *deleting* an employee from the *Employee* table should *also* result in all corresponding tuples from the *Works-on* relation to be *deleted*).
3. *Easier application development* and better *maintainability*:
  - Allows for the ICs to be **centrally managed** by the DBMS:
  - We do ***not*** have to ***worry*** about ***how*** integrity constraints are ***enforced/implemented***.



# Integrity Constraints in a Database

- › As previously indicated, ICs are specified as part of the database schema design
  - Note: The database designer is still *responsible* for ensuring that the integrity constraints *do not contradict* each other!
    - *The process of detection of IC conflicts may be automated, but this may introduce **unacceptable computational overhead**.*
- › Note: ICs are checked when the mentioned parts of the database are modified.
  - We can however specify when ICs should be *checked*: e.g., right *after* a SQL statement or at the *end* of a 'transaction'
    - *Transaction*: a **group of statements**, i.e., SQL statements, that are executed as *one single undivided* operation, i.e., **atomically**
- › What are the possible *ways to react* if an IC is *violated*?
  1. *Reject* the database operation – don't execute it.
  2. *Abort* the whole transaction – *rollback all* operations which are part of the transaction
  3. Execution of "*maintenance*" operations to make DB legal again

# An Informal Introduction to Transactions

Definition of a transaction: *A group of statements that are executed atomically*

BEGIN;

A group of SQL statements;

COMMIT;

› Note: An *SQL statement* starts with predefined keywords and ends with a semicolon `;`. These usually start with the keywords:

- **SELECT** - extracts data from a database
- **UPDATE** - updates data in a database
- **DELETE** - deletes data from a database
- **INSERT INTO** - inserts new data into a database
- **CREATE TABLE** - creates a new table
- **ALTER TABLE** - modifies a table
- **DROP TABLE** - deletes a table

› Consider an empty table

R(id: integer, name: varchar(8))

- What will be the result of the *following transaction*?

**BEGIN;**

**INSERT INTO R VALUES(1, 'Adam');**

**INSERT INTO R VALUES(1, 'Smith');**

**COMMIT;**

- Can you explain *why*?

## › *Two Types:*

### 1. *Static Integrity Constraints*

They describe conditions that *every legal instance* of a database *must satisfy*: Static integrity constraints are state **independent**.

- Inserts / deletes / updates which *violate* ICs are disallowed
- **Four types** of static integrity constraints:
  - *Domain Constraints*
  - *Key Constraints & Referential Integrity*
  - *Semantic Integrity Constraints*
  - *Assertions*

### 2. *Dynamic Integrity Constraints*

They are predicates on *database state changes* which capture constraints over *two or more states*. Therefore, dynamic integrity constraints are state **dependent**.

- *Triggers*

## › Overview of Integrity Constraints

- **Static Integrity Constraints**
  - Domain Constraints
  - Key / Referential Constraints
  - Semantic Integrity Constraints
  - Assertions
- **Dynamic Integrity Constraints**
  - Triggers

- › **Fields** must be of the *right data domain*
  - Always enforced for values that are *inserted* in the database
  - Queries are tested to ensure that the *comparisons* make *sense*.
- › Put simply, each attribute needs to have a *data type*
- › SQL DDL allows *domains* of attributes to be further *restricted* in the **CREATE TABLE** statement, using the following clauses:
  - **DEFAULT** *default-value*  
*default value* for an attribute if its value is *omitted* in an insert statement.
  - **NOT NULL**  
attribute is *not allowed* to become NULL
  - **NULL**  
the value of an attribute *may be* NULL (which is the default)

## Example of Domain Constraints

**CREATE TABLE** Student

```
(  
    sid      INTEGER      NOT NULL,  
    name     VARCHAR(20) NOT NULL,  
    semester INTEGER      DEFAULT 1,  
    birthday DATE         NULL,  
    country  VARCHAR(20)  
);
```

Example:

**INSERT INTO** Student(sid,name) **VALUES** (123,'Peter');

Student				
sid	name	semester	birthday	country
123	Peter	1	null	null

- › *Limit the allowed values* for an attribute by specifying *extra conditions* using an *in-line check* constraint

att-name sql-data-type **CHECK**(condition)

- › Examples:

- Grade can *only* be 'F','P','C','D', or 'H'

grade **CHAR**(1) **CHECK**( grade **IN** ('F','P','C','D','H'))

- Age *must* be positive

age **INTEGER CHECK**( age  $\geq 0$  )

- › *New domains* can be created from *existing data domains*, with their own defaults and restrictions

**CREATE DOMAIN** domain-name sql-data-type ...

- Example:

```
CREATE DOMAIN Grade CHAR(1) DEFAULT 'P' CHECK(VALUE IN  
('F','P','C','D','H'))
```


*Which is equivalent to:*

```
CREATE TABLE Student (  
  sid      INTEGER      NOT NULL,  
  name     VARCHAR(20)  NOT NULL,  
  grade    Grade,  
  birthday DATE);
```

```
CREATE TABLE Student (  
  sid      INTEGER      NOT NULL,  
  name     VARCHAR(20)  NOT NULL,  
  grade    CHAR (1) DEFAULT 'P' CHECK (grade IN ('F','P','C','D','H')),  
  birthday DATE);
```




- › In SQL, we specify *key* constraints using the **PRIMARY KEY** and **UNIQUE** clauses:



```
CREATE TABLE Student
(  
  sid    INTEGER PRIMARY KEY,  
  name  VARCHAR(20)  
);
```

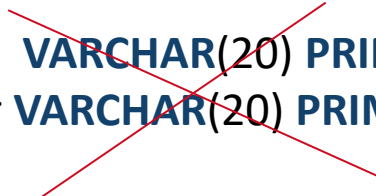
- › A **primary key** is automatically **UNIQUE** and **NOT NULL**
  - A relation can have **multiple** *candidate (unique)* keys, but **only** one *primary* key

- › Composite keys: a key consisting of multiple attributes
  - Must be specified in a separate clause



```
CREATE TABLE Boat
(
  name VARCHAR(20),
  colour VARCHAR(20),
  PRIMARY KEY (name, colour)
);
```

```
CREATE TABLE Boat
(
  name VARCHAR(20) PRIMARY KEY,
  colour VARCHAR(20) PRIMARY KEY
);
```



- › **Foreign key:** set of attributes in a *referring* relation that is used to `refer' to a tuple in a *parent/referred* relation.
    - Must refer to a **candidate key** of the parent (i.e., *referred*) relation
  
  - › **Definition of Referential Integrity:** For each tuple in the *referring* relation (called *child* or *dependent table*) whose foreign key value is **a**, there **must** be a tuple in the *referred* relation (called *parent table*) **whose value of the referred attribute is also a**
- e.g. Enrolled(sid: integer, ucode: char(8), semester: char(2))  
Student(sid: integer, name: varchar(20), age: integer, country: char(3))  
*sid* is a foreign key in *Enrolled* referring to the table *Student*:
- If all *foreign key constraints* are enforced, then referential integrity is *achieved*, i.e., *no dangling* references

- Express the constraint that *only students listed* in the Student relation should be allowed to *enrol* in Units of study.

```
CREATE TABLE Enrolled
( sid INTEGER, uos CHAR(8), grade VARCHAR(2),
  PRIMARY KEY (sid,uos),
  FOREIGN KEY (sid) REFERENCES Student,
  FOREIGN KEY (uos) REFERENCES Unitofstudy
);
```

**Student**

<u>sid</u>	name	age	country
53666	Jones	19	AUS
53650	Smith	21	AUS
54541	Ha Tschi	20	CHN
54672	Loman	20	AUS

**Enrolled**

<u>sid</u>	<u>uos</u>	grade
53666	COMP5138	CR
53666	INFO4990	CR
53650	COMP5138	P
53666	SOFT4200	D
<del>54221</del>	<del>INFO4990</del>	<del>F</del>

**Dependent Table**

??? Dangling reference

**Parent Table**

# Enforcing Referential Integrity in SQL

- › SQL-92 and SQL-99 support the following options on *deletes* and *updates* on the *parent* relation.
  - Default is **NO ACTION** (delete/update is rejected)
  - **CASCADE** (also delete/update *all tuples* that *refer* to deleted/updated tuple)
  - **SET NULL / SET DEFAULT** (also set *all* tuples that refer to deleted/updated tuple to **NULL/DEFAULT** values)

**CREATE TABLE** Enrolled

```
(  
  -- the sid field default value is 12345  
  sid CHAR(5) DEFAULT 12345,  
  uos CHAR(8),  
  grade VARCHAR(2),
```

**PRIMARY KEY** (sid,uos),

**FOREIGN KEY** (sid) **REFERENCES** Student  
**ON DELETE CASCADE**

-- the *on delete cascade* conveys  
that an *enrolled row* should be  
deleted when the student with sid  
that it refers to is deleted

**ON UPDATE SET DEFAULT**

-- the *on update set default*  
will update the value of sid to a default value  
that is specified as the default  
in this *enrolled* schema definition

);

# Table Constraints: Semantic Integrity Constraints

## › Examples:

- “Total marks are between 0 and 100”
- “Only lecturers of a course can give marks for that course.”

## › Use SQL **CHECK** constraints, *in-line* like before, or as *separate named* constraints:

**CHECK** ( semantic-condition )

# Example of Semantic Integrity Constraints

**CREATE TABLE** Assessment

```
(  
    sid    INTEGER    REFERENCES Student,  
    uos    VARCHAR(8) REFERENCES UnitOfStudy,  
    mark   INTEGER,  
    CHECK (mark BETWEEN 0 AND 100)  
);
```

› The **CONSTRAINT** clause can be used to *name* any integrity constraints

› Example:

```
CREATE TABLE Enrolled
(
  sid    INTEGER,
  uos    VARCHAR(8),
  grade  VARCHAR(2),
  CONSTRAINT FK_sid_enrolled FOREIGN KEY (sid)
                        REFERENCES Student
                        ON DELETE CASCADE,
  CONSTRAINT FK_cid_enrolled FOREIGN KEY (uos)
                        REFERENCES UnitOfStudy
                        ON DELETE CASCADE,
  CONSTRAINT CK_grade_enrolled CHECK(grade IN ('F',...)),
  CONSTRAINT PK_enrolled PRIMARY KEY (sid,uos)
);
```



# Timing of an Integrity Constraint Check

› Any constraint - *domain, key, foreign-key, or semantic* - may be declared as:

- **NOT DEFERRABLE**

**The default.** It means that every time a database modification occurs to tuples that a DBMS sees as being related, the constraint is ***checked immediately***.

- **DEFERRABLE**

Gives the option to ***wait until a transaction is complete*** and *then check* the constraint.

- **INITIALLY DEFERRED** wait until transaction ends  
however, can *dynamically* change this *later* in the transaction by

**SET CONSTRAINTS** name **IMMEDIATE**

- **INITIALLY IMMEDIATE** check immediate,  
however, can *dynamically* change this *later* in the transaction by

**SET CONSTRAINTS** name **DEFERRED**

# Example of Deferring Constraint Checking

```
CREATE TABLE UnitOfStudy
```

```
(
```

```
  uos_code      VARCHAR(8),
```

```
  title         VARCHAR(20),
```

```
  lecturer_id   INTEGER,
```

```
  credit_points INTEGER,
```

```
  CONSTRAINT UoS_PK PRIMARY KEY (uos_code),
```

```
  CONSTRAINT UoS_FK FOREIGN KEY (lecturer_id)
```

```
    REFERENCES Lecturer DEFERRABLE INITIALLY DEFERRED
```

```
);
```

- › Allows us to *insert a new unit of study* referencing a lecturer *who is not present at the time of adding a unit of study*, but *who will be added later in the same transaction*.
- › Behaviour can be *dynamically changed* within a transaction with the SQL statement:

```
SET CONSTRAINTS UoS_FK IMMEDIATE;
```

## Add/Modify/Remove Integrity Constraints\*

- › Integrity constraints can be *added*, *modified* (*only* applicable to *domain* constraints), and *removed* from an existing schema using **ALTER TABLE** statement

**ALTER TABLE** table-name constraint-modification

- › where *constraint-modification* is one of the following:

**ADD CONSTRAINT** constraint-name new-constraint

**DROP CONSTRAINT** constraint-name

**RENAME CONSTRAINT** old-name **TO** new-name

**ALTER COLUMN** attribute-name domain-constraint

- › Example (PostgreSQL syntax):

**ALTER TABLE** Enrolled **ALTER COLUMN** grade **TYPE VARCHAR(3)**,  
**ALTER COLUMN** mark **SET NOT NULL**;

- › What happens if the existing data in a table does not fulfil a newly added constraint?

The constraint **doesn't get created!**

e.g. "SQL Error: ORA-02296: cannot enable (USER.) - null values found"

Short break:

Let us have some fun again...



THE UNIVERSITY OF  
SYDNEY

# Schema-based Constraints: Assertions

- › The integrity constraints seen have so far been defined *within a single table*.
- › Some constraints *cannot* be expressed using *only domain constraints* or *referential integrity constraints*; for example:
  - “Every school must have at least five courses offered every semester” – must be expressed as an **assertion**
- › Need for ***more general integrity constraints*** that apply to the ***database: schema-based constraints***
  - Integrity constraints that span *several tables*
- › Definition of an **Assertion**: a ***predicate*** expressing a condition that we wish the ***database*** to ***always satisfy***.
- › SQL-92 syntax: **CREATE ASSERTION** assertion-name **CHECK** (condition)
  - When an *assertion* is made, the system *tests* it for its **validity** (i.e., it must *evaluate* to **true**) on any update which may potentially violate it
  - Note: This testing may introduce a *significant amount of overhead*; hence assertions should be used with great care.

- › For a **sailing club** to be categorized as ***small***, we require that the **sum** of the ***number of boats*** and ***number of sailors***, ***be less than 10*** at all times.

```
CREATE TABLE Sailors (  
  sid  INTEGER,  
  sname CHAR(10),  
  rating INTEGER,  
  PRIMARY KEY (sid),  
  CHECK (rating >=1 AND rating <=10),  
  CHECK ((SELECT COUNT(s.sid) FROM Sailors s)  
    + (SELECT COUNT(b.bid) FROM Boats b) < 10))  
);
```

Solution:

```
CREATE ASSERTION smallclub CHECK  
((SELECT COUNT(s.sid) FROM Sailors s)  
  + (SELECT COUNT(b.bid) FROM Boats b) < 10) );
```

What is *wrong* with the above SQL statement?

The check constraint would be checked only if there is an *update* in the table *Sailors*!

We could keep *adding* boats to be 10 or more and still having *no* constraint *violations*!

- › Note that *assertions of the form “for all X, P(X)”* are *not supported* by SQL and are typically *time-consuming*. Also, note that in *first order logic*:

$$\forall x P(x) \equiv \neg \exists x (\neg P(x))$$

( $\forall$  is read “*for all*”)

( $\exists$  is read “*there exists at least one*”)

( $\equiv$  means “equivalent”,  $\neg$  means “not”)

- ›  $\forall x P(x)$  reads like: For all values of x the condition P is true.
- ›  $\neg \exists x (\neg P(x))$  reads like: There does not exist any value x for which the opposite of condition P is true.



## Using General Assertions

- › How do we *declare* these types of assertions?
  - Write a query to ***select*** those tuples that ***violate*** the condition  **$P(x)$** , i.e., such that  **$\neg P(x)$**  is true.
  - Then use the **NOT EXISTS** clause, i.e.,  **$\neg \exists x$** . This ensures that the **assertion** would return ***true*** whenever the ***returned set*** is ***empty***!
    - The query result ***must therefore be empty*** if the ***assertion is to be satisfied***.
    - If the query result ***is not empty***, the assertion has been ***violated***!



Example: Assume we have *four* relations: ***Loan***, ***Borrower***, ***Depositor***, and ***Account***.

Define an **assertion** that states that *each loan has at least one borrower who maintains an account with a minimum balance of \$1000*. Assume ***x*** refers to tuples in ***Loan*** and ***y*** refers to tuples in ***Borrower***:

*This query can be written as follows:*

$$\forall x (\exists y P(x,y))$$

Another way to state this assertion is: *There are **no** loans that have **no** borrowers maintaining an account with a minimum balance of \$1000.*

The above query can be written as

$$\neg \exists x (\neg \exists y (P(x,y)))$$

where  $P(x,y) = (y.balance \geq 1000 \wedge x.loan\_number = y.loan\_number)$  where the  $\wedge$  symbol is read as “and”.

## Using General Assertions: Example

- › Using the previous formula:  $\neg\exists x (\neg\exists y (P(x,y)))$
- › **CREATE ASSERTION** balance\_constraint **CHECK** (*NOT EXISTS*
  - (**SELECT** \* **FROM** Loan **WHERE** *NOT EXISTS*  
(**SELECT** Borrower.loan\_number, Borrower.balance **FROM** Borrower **JOIN** Depositor  
ON Borrower.customer\_name = Depositor.customer\_name  
**JOIN** Account ON Depositor.account\_number = Account.account\_number  
**WHERE**  
Loan.loan\_number = Borrower.loan\_number  
**AND**  
Account.balance >= 1000))

$\neg\exists x$

$\neg\exists y$

$P(x,y)$

## Using General Assertions: Example

Given the previous assertion: *Each **loan** has at least one **borrower** who maintains an **account** that has a **minimum balance** of \$1000.* Assume we have the schema instantiation below.

Assume we wish to **update** the **Account** tuple (1221, 1200) with a **withdrawal** of \$700. If *successful*, this would take the balance down to **\$500**.

### Borrower

customer_name	address	loan_number
Jane	2, Happy Street	102
John	45, Glad Street	101

### Account

account_number	balance
2342	1500
1221	<del>1200</del>
	500?

### Depositor

account_number	customer_name
2342	Jane
1221	John

### Loan

loan_number	customer_name
101	John
102	Jane

The assertion is **checked** when the **update of (1221, 500)** in the table **Account** is attempted.

- Note there that the account number **(1221)** is associated with the borrower **(John)**.
- In this example, **loan\_number (101)** is associated with *only one* borrower **(John)**.
- If we allow the update to proceed, **loan\_number (101)** associated with the borrower **(John)** with the account **(1221)** would have a resulting **balance of (500)** which is **less than \$1000**, causing a **violation** of the assertion! In this case, the **update is rejected**.

## Using General Assertions: Example

- **Question:** what would the result of the update be if the *loan number of Jane* was **101**, i.e., it is a *joint loan* with John?
- **Answer:** the update would, in this case, be accepted!

**Borrower**

customer_name	address	loan_number
Jane	2, Happy Street	101
John	45, Glad Street	101

**Account**

account_number	balance
2342	1500
1221	<del>1200</del> 500

**Depositor**

account_number	customer_name
2342	Jane
1221	John

**Loan**

loan_number	customer_name
101	John
101	Jane

## Using General Assertions: Another Example\*

*The **sum of loan amounts** for each branch **must be less than or equal** to the **sum of all account balances** at the corresponding **branch**.*

Another way to state this assertion is: *There are **no** branches that have the **sum of loans greater** than the **sum of account balances** at that **branch**.*

```
> CREATE ASSERTION sum_constraint CHECK (NOT EXISTS  
  (SELECT * FROM branch  
    WHERE (SELECT sum(amount) FROM loan  
           WHERE loan.branch_name = branch.branch_name)  
           >  
           (SELECT sum(amount) FROM account  
           WHERE account.branch_name = branch.branch_name))))
```

Although **ASSERTION** is in the SQL standard, only a few DBMSs support it. **CHECK** are commonly used as a workaround approach.

*Assertion is still an effective design tool to use as the first instance to express constraints at the schema level.*

PostgreSQL **does not** support **ASSERTION** <https://www.postgresql.org/docs/9.2/unsupported-features-sql-standard.html>

PostgreSQL **does not** support **subquery in CHECK** <https://www.postgresql.org/docs/9.1/sql-createtable.html>

## › Overview of Integrity Constraints

### › Static Integrity Constraints

- Domain Constraints
- Key / Referential Constraints
- Semantic Integrity Constraints
- Assertions

### › Dynamic Integrity Constraints

- Triggers



- › A **trigger** is a piece of code that is executed *automatically* if some specified *modifications* occur to the database AND a *certain condition* holds true.
  
- › A trigger specification consists of *three* parts:  
    **ON event IF condition THEN action**
  - **Event**           ( what activates the trigger)
  - **Condition**      ( test the condition's truth to determine whether to execute an action)
  - **Action**          ( what happens if the condition is true)
  - Acronym is **ECA**
  - *Not all triggers have a condition, i.e., upon an event an action may be fired.*
  - Triggers are the basis of a specialized type of databases called **Active Databases**.

- › Note that Assertions ***do not modify*** the data:
  - Need a more powerful mechanism to *check conditions* and *modify* the data in a database.
- › Examples of use of triggers:
  - **Constraint maintenance**
    - Triggers can be used to *maintain* foreign-key and semantic constraints; commonly used with ON DELETE and ON UPDATE
  - **Business rules**
    - Dynamic business rules can be encoded as triggers. E.g., rules to adjust discounts based on stock availability
  - **Monitoring**
    - E.g., to react on the insertion of some sensor readings.
  - **Auditing/Compliance requirements**
    - E.g., government reporting regulations

## Trigger Example (SQL:1999)

Example: If the sum of marks of all current assessments for a student is greater than or equal to 50, enter the grade “P”.

**Event:** *new or updated assessment*

**Condition:** *check whether the sum of marks is greater than or equal to 50*

**Action:** *Enter grade “P” if the **condition** evaluates to true*

```
CREATE TRIGGER gradeEntry
  AFTER INSERT OR UPDATE ON Assessment
  BEGIN
    UPDATE Enrolled E
      SET grade='P'
    WHERE ( SELECT SUM(mark)
             FROM Assessment A
            WHERE A.sid=E.sid AND
                  A.uos=E.uosCode ) >= 50;
  END;
```

Event

Action

Condition

- › Triggering event can be **INSERT**, **DELETE** or **UPDATE**
- › Triggers on **update** can be **restricted** to **specific attributes**

**CREATE TRIGGER** overdraft-trigger

**AFTER UPDATE OF** *balance* **ON** Account

- › Values of attributes **before** and **after** an update can be *referenced*:
  - **REFERENCING OLD ROW AS** name: for *deletes* and *updates*
  - **REFERENCING NEW ROW AS** name: for *inserts* and *updates*
  - In PostgreSQL: separate **OLD** and **NEW** variable automatically generated with a trigger function (PL/pgsql).

## › Granularity

- **Row-level trigger:** A **row-level trigger** is fired for **each row** that needs to be updated.
  - **Statement-level trigger:** A **statement trigger** fires **once per triggering event** and *regardless of how many rows are modified* by the insert, update, or delete event.
- › *Statement-level* trigger is usually more *efficient* when dealing with SQL statements that update *many rows as a group*:
- if several rows in a table are updated, then a *statement-level* trigger would, in this case, be executed **only once**.

## Example of Row Level Trigger

Example: Assume that for auditing purposes, we define a **trigger** that ensures that any time a row is *inserted or updated* in the *employee* table, *the current user name of the person that made the modification and the time, are timestamped into the row*. It also checks that an *employee's name is provided* and that *the salary is a positive value*.

```
CREATE TABLE emp (  
    empname text,  
    salary integer,  
    last_date timestamp,  
    last_user text  
);  
  
CREATE FUNCTION emp_stamp() RETURNS trigger AS $emp_stamp$  
BEGIN  
    -- Check that empname and salary are given  
    IF NEW.empname IS NULL THEN  
        RAISE EXCEPTION 'empname cannot be null';  
    END IF;  
    IF NEW.salary IS NULL THEN  
        RAISE EXCEPTION '% cannot have null salary', NEW.empname;  
    END IF;  
  
    -- check if salary is negative  
    IF NEW.salary < 0 THEN  
        RAISE EXCEPTION '% cannot have a negative salary', NEW.empname;  
    END IF;  
  
    -- Remember who changed the payroll when  
    NEW.last_date := current_timestamp;  
    NEW.last_user := current_user;  
    RETURN NEW;  
END;  
$emp_stamp$ LANGUAGE plpgsql;  
CREATE TRIGGER emp_stamp BEFORE INSERT OR UPDATE ON emp  
FOR EACH ROW EXECUTE PROCEDURE emp_stamp();
```

- › Instead of executing a separate action for each affected row, a **single action** can be executed **for all rows** affected by a **statement**
  - Use **FOR EACH STATEMENT** (this is the **default**) instead of **FOR EACH ROW**
- › Statement-level triggers are more efficient when dealing with SQL statements that *update many rows as a group*.

Example: Assume the average salary is *timestamped* in a *dedicated relation* every time new employees are inserted or employee records are updated

```
CREATE FUNCTION Salary_Average() RETURNS trigger AS $$  
    BEGIN  
        INSERT INTO salaryaverages(datestamp,average) VALUES  
        (CURRENT_DATE, (SELECT AVG(salary) FROM Employee));  
        RETURN null;  
    END;  
$$ LANGUAGE plpgsql;  
CREATE TRIGGER RecordNewAverage  
AFTER UPDATE OF Salary or INSERT ON Employee  
FOR EACH STATEMENT  
    EXECUTE PROCEDURE Salary_Average();
```



```
CREATE [OR REPLACE] TRIGGER trigger-name
```

```
    ( BEFORE  
      AFTER  
      INSTEAD OF ) ( INSERT  
                     DELETE  
                     UPDATE OF attr ) ON table-name
```

```
    REFERENCEING    ( OLD  
                     NEW ) TABLE AS variable-name --optional
```

```
    FOR EACH ROW
```

```
    WHEN (condition)
```

```
    DECLARE
```

```
        <local variable declarations>
```

```
    BEGIN
```

```
        <PL/SQL block>
```

```
    END;
```

-- optional; otherwise, a statement trigger  
-- optional

In PostgreSQL, this  
is replaced by a  
trigger procedure

- › Things to consider when deciding to use a **row** or **statement** level trigger:

## **Update Cost**

- How many rows are updated?
- How often is a ***row-level*** trigger executed?
- How often is a ***statement-level*** trigger executed?

# Statement vs Row Level Trigger: A Recap

## Row Level Triggers

vs.

## Statement Level Triggers

Row level triggers executes once for each and every row that is updated/inserted/deleted.

Statement level triggers executes only once for each single SQL statement.

Mostly used for data auditing purpose.

Used for bulk modifications performed on the table.

“FOR EACH ROW” clause is present in CREATE TRIGGER command.

“FOR EACH ROW” clause is omitted in CREATE TRIGGER command.

Example: If 1500 rows are to be inserted into a table, the row level trigger would execute 1500 times.

Example: If 1500 rows are to be inserted into a table in one single statement, the statement level trigger would execute only once.

- › Use BEFORE triggers
  - Usually for *checking integrity constraints*
- › Use AFTER triggers
  - Usually for *integrity maintenance and update propagation*

- › Capture Integrity Constraints in an SQL Schema
  - Including key constraints, referential integrity, domain constraints and semantic constraints
- › Formulate complex semantic constraints using Assertions
- › Know when to use Assertions and CHECK constraints
- › Know the semantic of deferring integrity constraints
- › Be able to formulate simple triggers
  - Know the difference between row-level & statement-level triggers

- › Ramakrishnan/Gehrke (3rd edition - the 'Cow' book)
  - **Sections 3.2-3.3 and Sections 5.7-5.9**
  - *Integrity constraints are covered in different parts of the SQL discussion; only brief on triggers*
- › Kifer/Bernstein/Lewis (2nd edition)
  - Sections 3.2.2-3.3 and Chapter 7
  - *Integrity constraints are covered as part of the relational model, but there is a good dedicated chapter (Chap 7) on triggers*
- › Ullman/Widom (3rd edition)
  - Chapter 7
  - *Has a complete chapter dedicated to both integrity constraints & triggers.*
- › Michael v. Mannino: "Database - Design, Application Development and Administration"
  - *Includes a good introduction to triggers.*

› Advanced SQL

- Nested Queries
- Aggregation & Grouping
- NULL Values

› Readings:

- **Ramakrishnan/Gehrke (Cow book), Chapter 5 & Chapter 4.2.5**
- Kifer/Bernstein/Lewis book, Chapter 5 & 3.2-3.3
- Ullman/Widom, Chapter 6

See you next time!



THE UNIVERSITY OF  
SYDNEY